

NAME

TTF2UFM – A True Type to PostScript Type 1 Font Converter

SYNOPSIS

```
ttf2ufm [-options] ttf font.ttf [Fontname]
```

or

```
ttf2ufm [-options] ttf font.ttf -
```

DESCRIPTION

Ttf2ufm is a font converter from the True Type format (and some other formats supported by the FreeType library as well) to the Adobe Type1 format.

The versions 3.0 and later got rather extensive post-processing algorithm that brings the converted fonts to the requirements of the Type1 standard, tries to correct the rounding errors introduced during conversions and some simple kinds of bugs that are typical for the public domain TTF fonts. It also generates the hints that enable much better rendering of fonts in small sizes that are typical for the computer displays. But everything has its price, and some of the optimizations may not work well for certain fonts. That's why the options were added to the converter, to control the performed optimizations.

OPTIONS

The first variant creates the file `Fontname.pfa` (or `Fontname.pfb` if the option '**-b**' was used) with the converted font and `Fontname.afm` with the font metrics, the second one prints the font or another file (if the option '**-G**' was used) on the standard output from where it can be immediately piped through some filter. If no `Fontname` is specified for the first variant, the name is generated from `ttffont` by replacing the `.ttf` filename suffix.

Most of the time no options are necessary (with a possible exception of '**-e**'). But if there are some troubles with the resulting font, they may be used to control the conversion. The **options** are:

- **-a** – Include all the glyphs from the source file into the converted file. If this option is not specified then only the glyphs that have been assigned some encoding are included, because the rest of glyphs would be inaccessible anyway and would only consume the disk space. But some applications are clever enough to change the encoding on the fly and thus use the other glyphs, in this case they could benefit from using this option. But there is a catch: the X11 library has rather low limit for the font size. Including more glyphs increases the file size and thus increases the chance of hitting this limit. See `app/X11/README` for the description of a patch to X11 which fixes this problem.
- **-b** – Encode the resulting font to produce a ready `.pfb` file.
- **-d suboptions** – Debugging options. The suboptions are:
 - a** – Print out the absolute coordinates of dots in outlines. Such a font can not be used by any program (that's why this option is incompatible with '**-e**') but it has proven to be a valuable debugging information.
 - r** – Do not reverse the direction of outlines. The TTF fonts have the standard direction of outlines opposite to the Type1 fonts. So they should be reversed during proper conversion. This option may be used for debugging or to handle a TTF font with wrong direction of outlines (possibly, converted in a broken way from a Type1 font). The first signs of the wrong direction are the letters like "P" or "B" without the unpainted "holes" inside.
- **-e** – Assemble the resulting font to produce a ready `.pfa` file.

[S.B.: Personally I don't think that this option is particularly useful. The same result may be achieved by piping the unassembled data through `t1asm`, the Type 1 assembler. And, anyways, it's good to have the `t1utils` package handy. But Mark and many users think that this functionality is good and it took not much time to add this option.]
- **-F** – Force the Unicode encoding: any type of MS encoding specified in the font is ignored and the font is treated like it has Unicode encoding. **WARNING:** this option is intended for buggy fonts which actually are in Unicode but are marked as something else. The effect on the other fonts is unpredictable.

- **-G suboptions** – File generation options. The suboptions may be lowercase or uppercase, the lowercase ones disable the generation of particular files, the corresponding uppercase suboptions enable the generation of the same kind of files. If the result of ttf2ufm is requested to be printed on the standard output, the last enabling suboption of **-G** determines which file will be written to the standard output and the rest of files will be discarded. For example, **-G A** will request the AFM file. The suboptions to disable/enable the generation of the files are:

f/F – The font file. Depending on the other options this file will have one of the suffixes `.t1a`, `.pfa` or `.pfb`. If the conversion result is requested on the standard output (`'-'` is used as the output file name) then the font file will also be written there by default, if not overwritten by another suboption of **-G**.
Default: enabled

a/A – The Adobe font metrics file (`.afm`). **Default: enabled**

e/E – The dvips encoding file (`.enc`). **Default: disabled**

- **-l language [+argument]** – Extract the fonts for the specified language from a multi-language Unicode font. If this option is not used the converter tries to guess the language by the values of the shell variable LANG. If it is not able to guess the language by LANG it tries all the languages in the order they are listed.

After the plus sign an optional argument for the language extractor may be specified. The format of the argument is absolutely up to the particular language converter. The primary purpose of the argument is to support selection of planes for the multi-plane Eastern encodings but it can also be used in any other way. The language extractor may decide to add the plane name in some form to the name of the resulting font. None of the currently supported languages make any use of the argument yet.

As of now the following languages are supported:

`latin1` – for all the languages using the Latin-1 encoding
`latin2` – for the Central European languages
`latin4` – for the Baltic languages
`latin5` – for the Turkish language
`cyrillic` – for the languages with Cyrillic alphabet
`russian` – historic synonym for `cyrillic`
`bulgarian` – historic synonym for `cyrillic`
`adobestd` – for the AdobeStandard encoding used by TeX
`plane+argument` – to select one plane from a multi-byte encoding

The argument of the "plane" language may be in one of three forms:

`plane+pid=<pid>,eid=<eid>`
`plane+pid=<pid>,eid=<eid>,<plane_number>`
`plane+<plane_number>`

Pid (TTF platform id) and eid (TTF encoding id) select a particular TTF encoding table in the original font. They are specified as decimal numbers. If this particular encoding table is not present in the font file then the conversion fails. The native ("ttf") front-end parser supports only pid=3 (Windows platform), the FreeType-based ("ft") front-end supports any platform. If pid/eid is not specified then the TTF encoding table is determined as usual: Unicode encoding if it's first or an 8-bit encoding if not (and for an 8-bit encoding the plane number is silently ignored). To prevent the converter from falling back to an 8-bit encoding, specify the Unicode pid/eid value explicitly.

Plane_number is a hexadecimal (if starts with "0x") or decimal number. It gives the values of upper bytes for which 256 characters will be selected. If not specified, defaults to 0. It is also used as a font

name suffix (the leading "0x" is not included into the suffix).

NOTE: You may notice that the language names are not uniform: some are the names of particular languages and some are names of encodings. This is because of the different approaches. The original idea was to implement a conversion from Unicode to the appropriate Windows encoding for a given language. And then use the translation tables to generate the fonts in whatever final encodings are needed. This would allow to pile together the Unicode fonts and the non-Unicode Windows fonts for that language and let the program to sort them out automatically. And then generate fonts in all the possible encodings for that language. An example of this approach is the Russian language support. But if there is no multiplicity of encodings used for some languages and if the non-Unicode fonts are not considered important by the users, another way would be simpler to implement: just provide only one table for extraction of the target encoding from Unicode and don't bother with the translation tables. The latin* "languages" are examples of this approach. If somebody feels that he needs the Type1 fonts both in Latin-* and Windows encodings he or she is absolutely welcome to submit the code to implement it.

WARNING: Some of the glyphs included into the AdobeStandard encoding are not included into the Unicode standard. The most typical examples of such glyphs are ligatures like 'fi', 'fl' etc. Because of this the font designers may place them at various places. The converter tries to do its best, if the glyphs have honest Adobe names and/or are placed at the same codes as in the Microsoft fonts they will be picked up. Otherwise a possible solution is to use the option '-L' with an external map.

- **-L file[+[pid=<pid>,eid=<eid>],[plane]]** – Extract the fonts for the specified language from a multi-language font using the map from this file. This is rather like the option '-I' but the encoding map is not compiled into the program, it's taken from that file, so it's easy to edit. Examples of such files are provided in maps/adobe-standard-encoding.map, CP1250.map. (**NOTE:** the 'standard encoding' map does not include all the glyphs of the AdobeStandard encoding, it's provided only as an example.) The description of the supported map formats is in the file maps/unicode-sample.map.

Likewise to '-I', an argument may be specified after the map file name. But in this case the argument has fixed meaning: it selects the original TTF encoding table (the syntax is the same as in '-I plane') and/or a plane of the map file. The plane name also gets added after dash to the font name. The plane is a concept used in the Eastern fonts with big number of glyphs: one TTF font gets divided into multiple Type1 fonts, each containing one plane of up to 256 glyphs. But with a little creativity this concept may be used for other purposes of combining multiple translation maps into one file. To extract multiple planes from a TTF font `ttf2ufm` must be run multiple times, each time with a different plane name specified.

The default original TTF encoding table used for the option '-L' is Unicode. The map files may include directives to specify different original TTF encodings. However if the pid/eid pair is specified with it overrides any original encoding specified in the map file.

- **-m type=value** – Set maximal or minimal limits of resources. These limits control the the font generation by limiting the resources that the font is permitted to require from the PostScript interpreter. The currently supported types of limits are:

h – the maximal hint stack depth for the substituted hints. The default value is 128, according to the limitation in X11. This seems to be the lowest (and thus the safest) widespread value. To display the hint stack depth required by each glyph in a .tla file use the script `scripts/cntstems.pl`.

- **-O suboptions** – Outline processing options. The suboptions may be lowercase or uppercase, the lowercase ones disable the features, the corresponding uppercase suboptions enable the same features. The suboptions to disable/enable features are:

b/B – Guessing of the ForceBold parameter. This parameter helps the Type1 engine to rasterize the bold fonts properly at small sizes. But the algorithm used to guess the proper value of this flag makes that guess based solely on the font name. In rare cases that may cause errors, in these cases you may want to disable this guessing. **Default: enabled**

h/H – Autogeneration of hints. The really complex outlines may confuse the algorithm, so theoretically it may be useful sometimes to disable them. Although up to now it seems that even bad hints are better

than no hints at all. **Default: enabled**

u/U – Hint substitution. Hint substitution is a technique permitting generation of more detailed hints for the rasterizer. It allows to use different sets of hints for different parts of a glyph and change these sets as necessary during rasterization (that's why “substituted”). So it should improve the quality of the fonts rendered at small sizes. But there are two catches: First, the X11 library has rather low limit for the font size. More detailed hints increase the file size and thus increase the chance of hitting this limit (that does not mean that you shall hit it but you may if your fonts are particularly big). This is especially probable for Unicode fonts converted with option **'-a'**, so you may want to use **'-a'** together with **'-Ou'**. See `app/X11/README` for the description of a patch to X11 which fixes this problem. Second, some rasterizers (again, X11 is the typical example) have a limitation for total number of hints used when drawing a glyph (also known as the hint stack depth). If that stack overflows the glyph is ignored. Starting from version 3.22 `ttf2ufm` uses algorithms to minimizing this depth, with the trade-off of slightly bigger font files. The glyphs which still exceed the limit set by option **'-mh'** have all the substituted hints removed and only base hints left. The algorithms seem to have been refined far enough to make the fonts with substituted hints look better than the fonts without them or at least the same. Still if the original fonts are not well-designed the detailed hinting may emphasize the defects of the design, such as non-even thickness of lines. So provided that you are not afraid of the X11 bug the best idea would be to generate a font with this feature and without it, then compare the results using the program `other/cmpf` (see the description in `other/README`) and decide which one looks better. **Default: enabled**

o/O – Space optimization of the outlines' code. This kind of optimization never hurts, and the only reason to disable this feature is for comparison of the generated fonts with the fonts generated by the previous versions of converter. Well, it `_almost_` never hurts. As it turned out there exist some brain-damaged printers which don't understand it. Actually this feature does not change the outlines at all. The Type 1 font manual provides a set of redundant operators that make font description shorter, such as `'10 hlineto'` instead of `'0 10 rlineto'` to describe a horizontal line. This feature enables use of these operators. **Default: enabled**

s/S – Smoothing of outlines. If the font is broken in some way (even the ones that are not easily noticeable), such smoothing may break it further. So disabling this feature is the first thing to be tried if some font looks odd. But with smoothing off the hint generation algorithms may not work properly too. **Default: enabled**

t/T – Auto-scaling to the 1000x1000 Type1 standard matrix. The TTF fonts are described in terms of an arbitrary matrix up to 4000x4000. The converted fonts must be scaled to conform to the Type1 standard. But the scaling introduces additional rounding errors, so it may be curious sometimes to look at the font in its original scale. **Default: enabled**

v/V – Do vectorization on the bitmap fonts. Functionally “vectorization” is the same thing as “autotracing”, a different word is used purely to differentiate it from the Autotrace library. It tries to produce nice smooth outlines from bitmaps. This feature is still a work in progress though the results are already mostly decent. **Default: disabled**

w/W – Glyphs' width corection. This option is designed to be used on broken fonts which specify too narrow widths for the letters. You can tell that a font can benefit from this option if you see that the characters are smashed together without any whitespace between them. This option causes the converter to set the character widths to the actual width of this character plus the width of a typical vertical stem. But on the other hand the well-designed fonts may have characters that look better if their widths are set slightly narrower. Such well-designed fonts will benefit from disabling this feature. You may want to convert a font with and without this feature, compare the results and select the better one. This feature may be used only on proportional fonts, it has no effect on the fixed-width fonts. **Default: disabled**

z/Z – Use the Autotrace library on the bitmap fonts. The results are horrible and **the use of this option is not recommended**. This option is present for experimental purposes. It may change or be removed in the future. The working tracing can be achieved with option **-OV**. **Default: disabled**

- **-p parser_name** – Use the specified front-end parser to read the font file. If this option is not used, ttf2ufm selects the parser automatically based on the suffix of the font file name, it uses the first parser in its list that supports this font type. Now two parsers are supported:

ttf – built-in parser for the ttf files (suffix .ttf)

bdf – built-in parser for the BDF files (suffix .bdf)

ft – parser based on the FreeType-2 library (suffixes .ttf, .otf, .pfa, .pfb)

The parser `ft` is **NOT** linked in by default. See `Makefile` for instructions how to enable it. We do not support this parser on Windows: probably it will work but nobody tried and nobody knows how to build it.

The conversion of the bitmap fonts (such as BDF) is simplistic yet, producing jagged outlines. When converting such fonts, it might be a good idea to turn off the hint substitution (using option **-Ou**) because the hints produced will be huge but not adding much to the quality of the fonts.

- **-u number** – Mark the font with this value as its UniqueID. The UniqueID is used by the printers with the hard disks to cache the rasterized characters and thus significantly speed-up the printing. Some of those printers just can't store the fonts without UniqueID on their disk. The problem is that the ID is supposed to be unique, as its name says. And there is no easy way to create a guaranteed unique ID. Adobe specifies the range 4000000–4999999 for private IDs but still it's difficult to guarantee the uniqueness within it. So if you don't really need the UniqueID don't use it, it's optional. Luckily there are a few millions of possible IDs, so the chances of collision are rather low. If instead of the number a special value 'A' is given then the converter generates the value of UniqueID automatically, as a hash of the font name. (NOTE: in the version 3.22 the algorithm for autogeneration of UniqueID was changed to fit the values into the Adobe-specified range. This means that if UniqueIDs were used then the printer's cache may need to be flushed before replacing the fonts converted by an old version with fonts converted by a newer version). A simple way to find if any of the fonts in a given directory have duplicated UniqueIDs is to use the command:

```
cat *.pf[ab] | grep UniqueID | sort | uniq -c | grep -v ' 1 '
```

Or if you use `scripts/convert` it will do that for you automatically plus it will also give the exact list of files with duplicate UIDs.

- **-v size** – Re-scale the font to get the size of a typical uppercase letter somewhere around the specified size. Actually, it re-scales the whole font to get the size of one language-dependent letter to be at least of the specified size. Now this letter is "A" in all the supported languages. The size is specified in the points of the Type 1 coordinate grids, the maximal value is 1000. This is an experimental option and should be used with caution. It tries to increase the visible font size for a given point size and thus make the font more readable. But if overused it may cause the fonts to look out of scale. As of now the interesting values of size for this option seem to be located mostly between 600 and 850. This re-scaling may be quite useful but needs more experience to understand the balance of its effects.
- **-W level** – Select the verbosity level of the warnings. Currently the levels from 0 to 4 are supported. Level 0 means no warnings at all, level 4 means all the possible warnings. The default level is 3. Other levels may be added in the future, so using the level number 99 is recommended to get all the possible warnings. Going below level 2 is not generally recommended because you may miss valuable information about the problems with the fonts being converted.
- **Obsolete option: -A** – Print the font metrics (.afm file) instead of the font on STDOUT. Use **-GA** instead.
- **Very obsolete option:**

The algorithm that implemented the forced fixed width had major flaws, so it was disabled. The code is still in the program and some day it will be refined and returned back. Meanwhile the option name **'-f'** was reused for another option. The old version was:

- **-f** – Don't try to force the fixed width of font. Normally the converter considers the fonts in which the

glyph width deviates by not more than 5% as buggy fixed width fonts and forces them to have really fixed width. If this is undesirable, it can be disabled by this option.

The `.pfa` font format supposes that the description of the characters is binary encoded and encrypted. This converter does not encode or encrypt the data by default, you have to specify the option `'-e'` or use the `tlasm` program to assemble (that means, encode and encrypt) the font program. The `tlasm` program that is included with the converter is actually a part of the `tlutils` package, rather old version of which may be obtained from

<http://ttf2ufm.sourceforge.net/tlutils.tar.gz>

Note that `tlasm` from the old version of that package won't work properly with the files generated by `ttf2ufm` version 3.20 and later. Please use `tlasm` packaged with `ttf2ufm` or from the new version `tlutils` instead. For a newer version of `tlutils` please look at

<http://www.lcdf.org/~eddi2two/type/>

EXAMPLES

So, the following command lines:

```
ttf2ufm -e ttffont.ttf t1font
```

```
ttf2ufm ttffont.ttf - | tlasm >t1font.pfa
```

represent two ways to get a working font. The benefit of the second form is that other filters may be applied to the font between the converter and assembler.

FILES

- TTF2UFM_LIBXDIR/tlasm
- TTF2UFM_SHAREDIR/*
- TTF2UFM_SHAREDIR/scripts/*
- TTF2UFM_SHAREDIR/other/*
- TTF2UFM_SHAREDIR/README
- TTF2UFM_SHAREDIR/FONTS

SEE ALSO

- *ttf2ufm_convert*(1)
- *ttf2ufm_x2gs*(1)
- *tlasm*(1)
- `ttf2ufm-announce@lists.sourceforge.net`

The mailing list with announcements about `ttf2ufm`. It is a moderated mailing with extremely low traffic. Everyone is encouraged to subscribe to keep in touch with the current status of project. To subscribe use the Web interface at <http://lists.sourceforge.net/mailman/listinfo/ttf2ufm-announce>. If you have only e-mail access to the Net then send a subscribe request to the development mailing list `ttf2ufm-devel@lists.sourceforge.net` and somebody will help you with subscription.

- `ttf2ufm-devel@lists.sourceforge.net`
- `ttf2ufm-users@lists.sourceforge.net`

The `ttf2ufm` mailing lists for development and users issues. They have not that much traffic either. To subscribe use the Web interface at <http://lists.sourceforge.net/mailman/listinfo/ttf2ufm-devel> and <http://lists.sourceforge.net/mailman/listinfo/ttf2ufm-users>. If you have only e-mail access to the Net then send a subscribe request to the development mailing list `ttf2ufm-devel@lists.sourceforge.net` and somebody will help you with subscription.

- <http://ttf2ufm.sourceforge.net>

The main page of the project.

<http://www.netspace.net.au/~mheath/ttf2ufm/>

The old main page of the project.

BUGS

It seems that many Eastern fonts use features of the TTF format that are not supported by the ttf2ufm's built-in front-end parser. Because of this for now we recommend using the FreeType-based parser (option '-p ft') with the "plane" language.

Troubleshooting and bug reports

Have problems with conversion of some font ? The converter dumps core ? Or your printer refuses to understand the converted fonts ? Or some characters are missing ? Or some characters look strange ?

Send the bug reports to the ttf2ufm development mailing list at ttf2ufm-devel@lists.sourceforge.net.

Try to collect more information about the problem and include it into the bug report. (Of course, even better if you would provide a ready fix, but just a detailed bug report is also good). Provide detailed information about your problem, this will speed up the response greatly. Don't just write "this font looks strange after conversion" but describe what's exactly wrong with it: for example, what characters look wrong and what exactly is wrong about their look. Providing a link to the original font file would be also a good idea. Try to do a little troubleshooting and report its result. This not only would help with the fix but may also give you a temporary work-around for the bug.

First, enable full warnings with option '-W99', save them to a file and read carefully. Sometimes the problem is with a not implemented feature which is reported in the warnings. Still, reporting about such problems may be a good idea: some features were missed to cut corners, in hope that no real font is using them. So a report about a font using such a feature may motivate someone to implement it. Of course, you may be the most motivated person: after all, you are the one wishing to convert that font. ;-) Seriously, the philosophy "scrath your own itch" seems to be the strongest moving force behind the Open Source software.

The next step is playing with the options. This serves a dual purpose: on one hand, it helps to localize the bug, on the other hand you may be able to get a working version of the font for the meantime while the bug is being fixed. The typical options to try out are: first '-Ou', if it does not help then '-Os', then '-Oh', then '-Oo'. They are described in a bit more detail above. Try them one by one and in combinations. See if with them the resulting fonts look better.

On some fonts ttf2ufm just crashes. Commonly that happens because the font being converted is highly defective (although sometimes the bug is in ttf2ufm itself). In any case it should not crash, so the reports about such cases will help to handle these defects properly in future.

We try to respond to the bug reports in a timely fashion but alas, this may not always be possible, especially if the problem is complex. This is a volunteer project and its resources are limited. Because of this we would appreciate bug reports as detailed as possible, and we would appreciate the ready fixes and contributions even more.

HISTORY

Based on ttf2pfa by Andrew Weeks, and help from Frank Siegart.

Modification by Mark Heath.

Further modification by Sergey Babkin.

The Type1 assembler by I. Lee Hetherington with modifications by Kai-Uwe Herbing.